

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import classification_report

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: data = pd.read_csv("E:/BE/Assignments/LP3/ML Assignments/1/uber.csv")
data.head()
```

```
Out[2]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085

```
In [3]: df = data.copy()
df.head()
```

```
Out[3]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount           200000 non-null  float64
3   pickup_datetime      200000 non-null  object
4   pickup_longitude      200000 non-null  float64
5   pickup_latitude       200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude      199999 non-null  float64
8   passenger_count       200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [5]: `df.describe()`

Out[5]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	2.000000e+05	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	2.771250e+07	11.359955	-72.527638	39.935885	-72.525292	39.935885
std	1.601382e+07	9.901776	11.437787	7.720539	13.117408	6.720539
min	1.000000e+00	-52.000000	-1340.648410	-74.015515	-3356.666300	-881.666300
25%	1.382535e+07	6.000000	-73.992065	40.734796	-73.991407	40.734796
50%	2.774550e+07	8.500000	-73.981823	40.752592	-73.980093	40.752592
75%	4.155530e+07	12.500000	-73.967154	40.767158	-73.963658	40.767158
max	5.542357e+07	499.000000	57.418457	1644.421482	1153.572603	872.666300

In [6]: `df.isnull().sum()`

Out[6]:

```
Unnamed: 0      0
key             0
fare_amount     0
pickup_datetime 0
pickup_longitude 0
pickup_latitude  0
dropoff_longitude 1
dropoff_latitude 1
passenger_count  0
dtype: int64
```

In [7]: `df = df.drop(['Unnamed: 0', 'key'], axis=1)`
`df.dropna(axis=0, inplace=True)`

```
In [8]: df.isnull().sum()
```

```
Out[8]: fare_amount          0
pickup_datetime            0
pickup_longitude           0
pickup_latitude            0
dropoff_longitude          0
dropoff_latitude           0
passenger_count            0
dtype: int64
```

Haversine Formula

Calculatin the distance between the pickup and drop co-ordinates using the Haversine formual for accuracy.

$$d = 2r \sin^{-1} \left(\sqrt{\sin^2 \left(\frac{\Phi_2 - \Phi_1}{2} \right) + \cos(\Phi_1) \cos(\Phi_2) \sin^2 \left(\frac{\lambda_2 - \lambda_1}{2} \right)} \right)$$

```
In [9]: def haversine (lon_1, lon_2, lat_1, lat_2):

    lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1, lat_2])

    diff_lon = lon_2 - lon_1
    diff_lat = lat_2 - lat_1

    km = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2 +
                                     np.cos(lat_1) * np.cos(lat_2) * np.sin(diff

    return km
```

```
In [10]: df['Distance'] = haversine(df['pickup_longitude'], df['dropoff_longitude'],
                                     df['pickup_latitude'], df['dropoff_latitude'])

df['Distance'] = df['Distance'].astype(float).round(2)
```


```
In [11]: df.head()
```

```
Out[11]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.7230
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.7500
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.7720
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.8030
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.7610

```
In [12]: plt.scatter(df['Distance'], df['fare_amount'])
plt.xlabel("Distance")
plt.ylabel("fare_amount")
```

```
Out[12]: Text(0, 0.5, 'fare_amount')
```



Outliers

We can get rid of the trips with very large distances that are outliers as well as trips with 0 distance.

```
In [13]: df.drop(df[df['Distance'] > 60].index, inplace = True)
df.drop(df[df['Distance'] == 0].index, inplace = True)
df.drop(df[df['Distance'] < 0].index, inplace = True)

df.drop(df[df['fare_amount'] == 0].index, inplace = True)
df.drop(df[df['fare_amount'] < 0].index, inplace = True)
```

```
In [14]: df.drop(df[df['Distance'] > 100].index, inplace = True)
df.drop(df[df['fare_amount'] > 100].index, inplace = True)
```

Also removing rows with non-plausible fare amounts and distance travelled


```
In [15]: df.drop(df[(df['fare_amount'] > 100) & (df['Distance'] < 1)].index, inplace = True)
df.drop(df[(df['fare_amount'] < 100) & (df['Distance'] > 100)].index, inplace = True)
```

```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 193436 entries, 0 to 199999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   fare_amount           193436 non-null float64
1   pickup_datetime       193436 non-null object
2   pickup_longitude      193436 non-null float64
3   pickup_latitude       193436 non-null float64
4   dropoff_longitude     193436 non-null float64
5   dropoff_latitude      193436 non-null float64
6   passenger_count       193436 non-null int64
7   Distance              193436 non-null float64
dtypes: float64(6), int64(1), object(1)
memory usage: 17.3+ MB
```

```
In [17]: plt.scatter(df['Distance'], df['fare_amount'])
plt.xlabel("Distance")
plt.ylabel("fare_amount")
```

```
Out[17]: Text(0, 0.5, 'fare_amount')
```



Separating the date and time into separate columns for more usability

```
In [18]: df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])

df['Year'] = df['pickup_datetime'].apply(lambda time: time.year)
df['Month'] = df['pickup_datetime'].apply(lambda time: time.month)
df['Day'] = df['pickup_datetime'].apply(lambda time: time.day)
df['Day of Week'] = df['pickup_datetime'].apply(lambda time: time.dayofweek)
df['Day of Week_num'] = df['pickup_datetime'].apply(lambda time: time.dayofweek)
df['Hour'] = df['pickup_datetime'].apply(lambda time: time.hour)

day_map = {0: 'Mon', 1: 'Tue', 2: 'Wed', 3: 'Thu', 4: 'Fri', 5: 'Sat', 6: 'Sun'}
df['Day of Week'] = df['Day of Week'].map(day_map)

df['counter'] = 1
```

Creating separate coumns for pickup and droppoff coordinates for more usability.

```
In [19]: df['pickup'] = df['pickup_latitude'].astype(str) + "," + df['pickup_longitude'].a
df['drop off'] = df['dropoff_latitude'].astype(str) + "," + df['dropoff_longitude']
```

```
In [20]: df.head()
```

Out[20]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.7232
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.7500
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.7726
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.8030
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.7612

Correlation

```
In [21]: corr = df.corr()

corr.style.background_gradient(cmap='BuGn')
```

Out[21]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.012292	-0.008891	0.010831	-0.00904
pickup_longitude	0.012292	1.000000	-0.949099	0.999885	-0.99397
pickup_latitude	-0.008891	-0.949099	1.000000	-0.949096	0.95476
dropoff_longitude	0.010831	0.999885	-0.949096	1.000000	-0.99396
dropoff_latitude	-0.009044	-0.993976	0.954760	-0.993964	1.00000
passenger_count	0.014409	0.009176	-0.009219	0.009164	-0.00926
Distance	0.895513	0.005356	0.003243	0.004464	-0.00225
Year	0.124050	0.013480	-0.013693	0.013373	-0.01436
Month	0.024850	-0.007497	0.007602	-0.007452	0.00798
Day	0.000277	0.019531	-0.019393	0.019555	-0.02011
Day of Week_num	0.004881	0.008243	-0.008924	0.008543	-0.00891
Hour	-0.020270	0.001835	-0.001821	0.000937	-0.00101
counter	nan	nan	nan	nan	na

There is some correlation between the distance and fare amount.
Implementing simple linear regression model using these two variables.

```
In [22]: X = df['Distance'].values.reshape(-1, 1)
y = df['fare_amount'].values.reshape(-1, 1)
```

In [23]: `from sklearn.preprocessing import StandardScaler`

```
std = StandardScaler()
y_std = std.fit_transform(y)
print(y_std)

x_std = std.fit_transform(X)
print(x_std)
```

```
[[-0.40638221]
 [-0.38489719]
 [ 0.17371326]
 ...
 [ 2.10736482]
 [ 0.3455934 ]
 [ 0.30262337]]
[[-0.46769936]
 [-0.24942881]
 [ 0.472543 ]
 ...
 [ 2.65804681]
 [ 0.05279195]
 [ 0.57887993]]
```

In [24]: `from sklearn.model_selection import train_test_split`

```
X_train, X_test, y_train, y_test = train_test_split(x_std, y_std, test_size=0.3,
```

Linear Regression Model

In [25]: `from sklearn.linear_model import LinearRegression`

```
l_reg = LinearRegression()
l_reg.fit(X_train, y_train)

print("Training set score: {:.2f}".format(l_reg.score(X_train, y_train)))
print("Test set score: {:.7f}".format(l_reg.score(X_test, y_test)))
```

```
Training set score: 0.80
Test set score: 0.8050688
```

In [27]: `y_pred = l_reg.predict(X_test)`


```
In [28]: from sklearn import metrics

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))

print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))

print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

Mean Absolute Error: 0.2438330049716194
Mean Squared Error: 0.1926995801043055
Root Mean Squared Error: 0.4389756030855308

```
In [29]: print(l_reg.intercept_)
print(l_reg.coef_)
```


[0.00029241]
[[0.89503692]]

Plotting the linear regression line against the training and test set side by side.

```
In [30]: plt.subplot(2, 2, 1)
plt.scatter(X_train, y_train, color = 'red')
plt.plot(X_train, l_reg.predict(X_train), color = "blue")
plt.title("Fare vs Distance (Training Set)")
plt.ylabel("fare_amount")
plt.xlabel("Distance")

plt.subplot(2, 2, 2)
plt.scatter(X_test, y_test, color = 'red')
plt.plot(X_train, l_reg.predict(X_train), color = "blue")
plt.ylabel("fare_amount")
plt.xlabel("Distance")
plt.title("Fare vs Distance (Test Set)")

plt.tight_layout()
plt.rcParams["figure.figsize"] = (32,22)
plt.show()
```



Random Forest Model

```
In [31]: from sklearn.ensemble import RandomForestRegressor

r_reg = RandomForestRegressor(n_estimators = 50, random_state = 0)

r_reg.fit(X_train, y_train)
```

Out[31]: RandomForestRegressor(n_estimators=50, random_state=0)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [32]: predictions = r_reg.predict(X_test)

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, predictions))

print('Mean Squared Error:', metrics.mean_squared_error(y_test, predictions))

print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

Mean Absolute Error: 0.24670010997672112
Mean Squared Error: 0.19753474247914912
Root Mean Squared Error: 0.4444488074898493

```
In [ ]:
```